

# Ubuntu, NGINX, uWSGI, and Flask Installation Guide on AWS

## Create Python and PIP alias

We need to first create an alias for python3 and pip3

```
sudo nano ~/.bashrc
```

Add these lines at the last of this file

```
alias python=python3
```

```
alias pip=pip3
```

## Create Extra RAM from SSD

```
sudo swapon --show
```

```
free -h
```

```
df -h
```

```
sudo fallocate -l 4G /swapfile
```

```
ls -lh /swapfile
```

```
sudo chmod 600 /swapfile
```

```
ls -lh /swapfile
```

```
sudo mkswap /swapfile
```

```
sudo swapon /swapfile
```

```
sudo swapon --show
```

```
free -h
```

```
sudo cp /etc/fstab /etc/fstab.bak
```

```
echo '/swapfile none swap sw 0 0' | sudo tee -a /etc/fstab
```

Type *Markdown* and LaTeX:  $\alpha^2$

## Update Ubuntu

Uninstall Anaconda if you have otherwise it will not work

You need to keep your system up-to-date

```
sudo apt update
```

```
sudo apt upgrade
```

```
sudo apt install python3-pip python3-dev build-essential libssl-dev libffi-dev python3-setuptools
```

```
pip3 install -U pip
```

```
pip3 install ktrain
```

```
pip3 install flask
```

## Install Python Virtual Environment

Install python virtual environment in order to isolate our new installation from previous working modules. This has to be done, incase anything breaks.

```
sudo apt install python3-venv
```

## Make a parent directory for our project

Here are few guidelines which we will follow

```
mkdir ~/mlmodel
```

```
cd ~/mlmodel
```

- project directory: mlmodel
- virtual environment: modelenv
- username: ubuntu
- server\_ip: EC2 Instance Public IP
- default http port: 80
- flask port: 5000

## Create a Virtual Environment

```
python3.6 -m venv modelenv
```

this will install local copy of Python packages into virtual environment

Now you need to activate virtual environment

```
source modelenv/bin/activate
```

Now you should see in terminal like this

```
(modelenv)ubuntu@host:~/mlmodel$
```

Congrats! Virtual Environment is activated

## Setting up Flask Application

Now we need to install necessary python packages inside modelenv

```
pip install -U pip
```

```
pip install wheel
```

```
pip install uwsgi flask
```

```
pip install ktrain
```

```
sudo ufw allow 5000
```

Now test your flask application `mlmodel.py`

Stop it after testing

## Setting up uWSGI

We need to create a `wsgi.py` file inside the `mlmodel` project dir

```
nano ~/mlmodel/wsgi.py
```

Inside `wsgi.py` file add below code

```
from mlmodel import app

if __name__ == "__main__":
    app.run()
```

## Testing uWSGI

specify the socket, so that it will be started on a publicly available interface, as well as the protocol, so that it will use HTTP instead of the uwsgi binary protocol. We'll use the same port number, `5000`, that we opened earlier

```
uwsgi --socket 0.0.0.0:5000 --protocol=http -w wsgi:app
```

now test server by visiting public ip address

```
http://public_ip:5000
```

Now deactivate the virtual environment

```
deactivate
```

Type *Markdown* and LaTeX:  $\alpha^2$

## Configuring uWSGI

We need to setup a robust system for deployment

```
nano ~/mlmodel/mlmodel.ini
```

we will start off with the `[uwsgi]` header so that uWSGI knows to apply the settings. We'll specify two things:

- the module itself, by referring to the `wsgi.py` file minus the extension,
- and the callable within the file, `app`:

```
[uwsgi]
module = wsgi:app

master = false
processes = 1
```

```
socket = mlmodel.sock
chmod-socket = 660
vacuum = true

die-on-term = true
```

The last thing we'll do is set the `die-on-term` option. This can help ensure that the init system and `uWSGI` have the same assumptions about what each process signal means.

You may have noticed that we did not specify a protocol like we did from the command line. That is because by default, `uWSGI` speaks using the `uwsgi` protocol, a fast binary protocol designed to communicate with other servers. `NGINX` can speak this protocol natively, so it's better to use this than to force communication by `HTTP`.

## Creating systemd Unit File

Creating a systemd unit file will allow Ubuntu's init system to automatically start `uWSGI` and serve the Flask application whenever the server boots.

```
sudo nano /etc/systemd/system/mlmodel.service
```

Write this below code

```
[Unit]
Description=uWSGI instance to serve mlmodel
After=network.target

[Service]
User=ubuntu
Group=www-data
WorkingDirectory=/home/ubuntu/mlmodel
Environment="PATH=/home/ubuntu/mlmodel/modelenv/bin"
ExecStart=/home/ubuntu/mlmodel/modelenv/bin/uwsgi --ini mlmodel.ini

[Install]
WantedBy=multi-user.target
```

## Start uWSGI services

We can now start the `uWSGI` service we created and enable it so that it starts at boot:

```
sudo systemctl start mlmodel
```

```
sudo systemctl enable mlmodel
```

```
sudo systemctl status mlmodel
```

Type *Markdown* and LaTeX:  $\alpha^2$

## Installing NGINX

```
sudo apt update
```

```
sudo apt install nginx
```

Nginx registers itself as a service with `ufw` upon installation, making it straightforward to allow `Nginx` access.

```
sudo ufw app list
```

We should allow only HTTPS but for time being we will use only HTTP. We will need additional SSL certificate to encrypt connection of https.

```
sudo ufw allow 'Nginx HTTP'
```

```
sudo ufw status
```

```
systemctl status nginx
```

Check your NGINX is working?

```
http://public_ip
```

Type *Markdown* and LaTeX:  $\alpha^2$

Managing NGINX process

```
sudo systemctl stop nginx
```

```
sudo systemctl start nginx
```

```
sudo systemctl restart nginx
```

```
sudo systemctl reload nginx
```

Type *Markdown* and LaTeX:  $\alpha^2$

## Configuring NGINX

```
sudo nano /etc/nginx/sites-available/mlmodel
```

```
server {
    listen 80;
    server_name mlmodel public_ip;

    location / {
        include uwsgi_params;
        uwsgi_pass unix:/home/ubuntu/mlmodel/mlmodel.sock;
    }
}
```

To enable the Nginx server block configuration you've just created, link the file to the `sites-enabled` directory:

```
sudo ln -s /etc/nginx/sites-available/mlmodel /etc/nginx/sites-enabled
```

```
sudo nginx -t
```

```
sudo systemctl restart nginx
```

```
sudo ufw delete allow 5000
```

```
sudo ufw allow 'Nginx Full'
```

```
http://public-ip
```

Type *Markdown* and LaTeX:  $\alpha^2$

## Errors checking

```
sudo less /var/log/nginx/error.log
```

```
sudo less /var/log/nginx/access.log
```

```
sudo journalctl -u nginx
```

```
sudo journalctl -u mlmodel
```

## API Test

```
In [2]: 1 import requests, json
```

```
In [3]: 1 url = "http://18.222.206.75/"
2
3 x = requests.get(url)
4 print(x)
5 print(x.text)
```

```
<Response [200]>
Congrats! Server is working
```

```
In [6]: 1 %%time
2 url = "http://18.222.206.75/get_prediction"
3 data = {'comment': 'this movie is horrible. return my money'}
4 data = json.dumps(data)
5
6 x = requests.post(url, data)
7 print(x.text)
```

```
{"result":"0"}
```

```
Wall time: 568 ms
```

```
In [ ]: 1
```